Towards Efficient AI: Techniques for Scalable and Resource-Constrained Models

Louis Sloot, Elizabeth Qiu, Maggie Chen
Mentors: Dr. Yifan (Evelyn) Gong, Joyce Li

Background

Deploying deep neural networks (DNNs) on resource-constrained devices is challenging due to high computational and storage demands, making efficiency critical.

- Energy consumption of AI has surged as model parameters grow exponentially
- Data centers account for ~3% of global carbon emissions annually

The increasing mainstream demand and popularity of large-scale AI models, including large language models (LLMs) and convolutional neural networks (CNNs), makes addressing these efficiency issues even more critical. **This study explores practical solutions to reduce model size, accelerate execution, and optimize training and inference processes.**

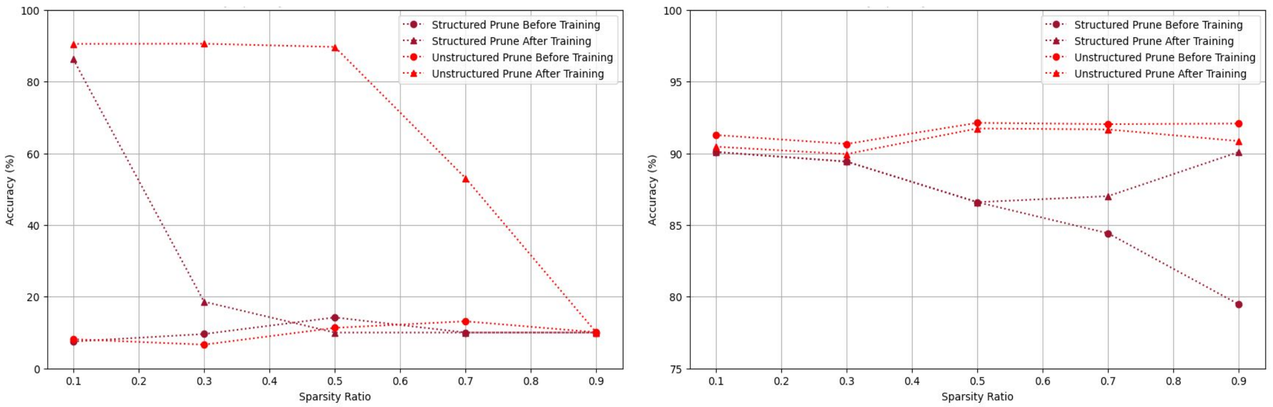

Figure 2 (*above*): Accuracy-Sparsity Ratio in Pruned Models: **Non-Retrained Models** (*Left*); **Retrained** (*Right*)

Table 1 (*below*): Multiply-Accumulate Operations (MACs) based on sparsity ratio

Sparsity Ratio	0.0	0.1	0.3	0.5	0.7	0.9
Predicted MACs (millions)	72.3	65.1	50.6	36.2	21.7	7.24

Takeaways

- Unstructured pruning more resistant to accuracy decrease, but theoretically harder to gain efficiency
- Structured pruning impacts accuracy; easier to leverage efficiency (smaller matrix dimensions)
- Overall, (small) portions of models can be pruned without critical loss in accuracy while making theoretical efficiency gains

Methods

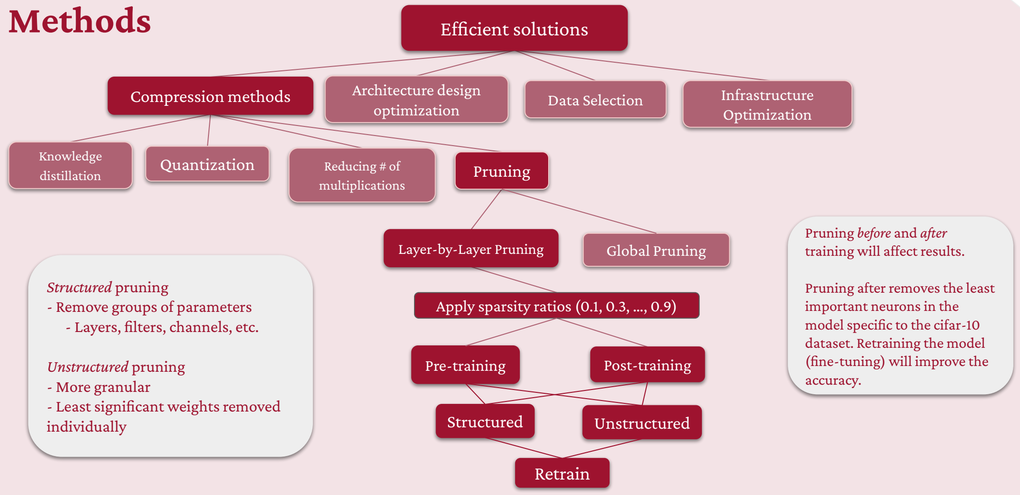

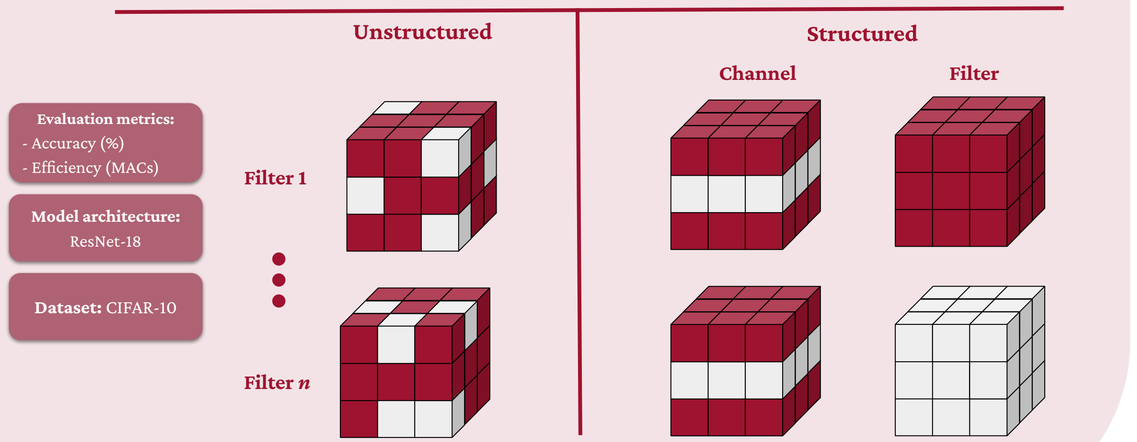

Figure 1: Pruning Methods

Future Actions

Gaining Efficiency from Zeroed Weights: Investigate, develop algorithm(s) that convert zeroed weights into lowered MACs and greater efficiency

Pruning During Training: Faster training times than pruning after training and higher accuracy than pruning before training

Sparsity Ratios: Varying sparsity ratios distributed across different layers accounts for some layers influencing results more than others

Experimenting with Different Norms: Currently using L1 norms to determine structured pruning, but other norms may yield different results

Carnegie Mellon University
School of Computer Science
University of Maryland
DEPARTMENT OF COMPUTER SCIENCE

Link to Codebase

Adobe